

Wine Quality Analysis and Prediction

```
In [1]: #importing the Libery
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
%matplotlib inline
warnings.filterwarnings('ignore')
```

Data Collection

```
In [2]: #loading dataset in pandas dataframe
data = pd.read_csv('/content/winequality.csv')
```

```
In [3]: #check first five rows of the dataset
data.head()
```

Out[3]:

	type	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphate
0	white	7.0	0.27	0.36	20.7	0.045	45.0	170.0	1.0010	3.00	0
1	white	6.3	0.30	0.34	1.6	0.049	14.0	132.0	0.9940	3.30	0
2	white	8.1	0.28	0.40	6.9	0.050	30.0	97.0	0.9951	3.26	0
3	white	7.2	0.23	0.32	8.5	0.058	47.0	186.0	0.9956	3.19	0
4	white	7.2	0.23	0.32	8.5	0.058	47.0	186.0	0.9956	3.19	0

└─

```
In [4]: #check last five rows of the dataset
data.tail()
```

Out[4]:

	type	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulfur
6492	red	6.2	0.600	0.08	2.0	0.090	32.0	44.0	0.99490	3.45	
6493	red	5.9	0.550	0.10	2.2	0.062	39.0	51.0	0.99512	3.52	
6494	red	6.3	0.510	0.13	2.3	0.076	29.0	40.0	0.99574	3.42	
6495	red	5.9	0.645	0.12	2.0	0.075	32.0	44.0	0.99547	3.57	
6496	red	6.0	0.310	0.47	3.6	0.067	18.0	42.0	0.99549	3.39	

```
In [5]: #check shape of the dataset
data.shape
```

Out[5]: (6497, 13)

```
In [6]: #check infomation of the dataset
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6497 entries, 0 to 6496
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   type                                  6497 non-null   object
1   fixed acidity                         6487 non-null   float64
2   volatile acidity                     6489 non-null   float64
3   citric acid                           6494 non-null   float64
4   residual sugar                       6495 non-null   float64
5   chlorides                           6495 non-null   float64
6   free sulfur dioxide                  6497 non-null   float64
7   total sulfur dioxide                 6497 non-null   float64
8   density                             6497 non-null   float64
9   pH                                   6488 non-null   float64
10  sulphates                           6493 non-null   float64
11  alcohol                             6497 non-null   float64
12  quality                             6497 non-null   int64
dtypes: float64(11), int64(1), object(1)
memory usage: 660.0+ KB
```

```
In [7]: #check columns of the dataset
data.columns
```

```
Out[7]: Index(['type', 'fixed acidity', 'volatile acidity', 'citric acid',
              'residual sugar', 'chlorides', 'free sulfur dioxide',
              'total sulfur dioxide', 'density', 'pH', 'sulphates', 'alco
              hol',
              'quality'],
              dtype='object')
```

```
In [8]: #check mathamatic describe
data.describe()
```

```
Out[8]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide
count	6487.000000	6489.000000	6494.000000	6495.000000	6495.000000	6497.000000	6497.000000
mean	7.216579	0.339691	0.318722	5.444326	0.056042	30.525319	118.411329
std	1.296750	0.164649	0.145265	4.758125	0.035036	17.749400	56.208919
min	3.800000	0.080000	0.000000	0.600000	0.009000	1.000000	6.000000
25%	6.400000	0.230000	0.250000	1.800000	0.038000	17.000000	71.000000
50%	7.000000	0.290000	0.310000	3.000000	0.047000	29.000000	118.000000
75%	7.700000	0.400000	0.390000	8.100000	0.065000	41.000000	156.000000
max	15.900000	1.580000	1.660000	65.800000	0.611000	289.000000	440.000000

```
In [10]: #check coreation of the dataset
data.corr()
```

Out[10]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	de
fixed acidity	1.000000	0.220172	0.323736	-0.112319	0.298421	-0.283317	-0.329747	0.45
volatile acidity	0.220172	1.000000	-0.378061	-0.196702	0.377167	-0.353230	-0.414928	0.27
citric acid	0.323736	-0.378061	1.000000	0.142486	0.039315	0.133437	0.195218	0.05
residual sugar	-0.112319	-0.196702	0.142486	1.000000	-0.128902	0.403439	0.495820	0.55
chlorides	0.298421	0.377167	0.039315	-0.128902	1.000000	-0.195042	-0.279580	0.36
free sulfur dioxide	-0.283317	-0.353230	0.133437	0.403439	-0.195042	1.000000	0.720934	0.02
total sulfur dioxide	-0.329747	-0.414928	0.195218	0.495820	-0.279580	0.720934	1.000000	0.03
density	0.459204	0.271193	0.096320	0.552498	0.362594	0.025717	0.032395	1.00
pH	-0.251814	0.260660	-0.328689	-0.267050	0.044806	-0.145191	-0.237687	0.01
sulphates	0.300380	0.225476	0.057613	-0.185745	0.395332	-0.188489	-0.275381	0.25
alcohol	-0.095603	-0.038248	-0.010433	-0.359706	-0.256861	-0.179838	-0.265740	-0.65
quality	-0.077031	-0.265953	0.085706	-0.036825	-0.200886	0.055463	-0.041385	-0.30

```
In [11]: #check missing value of the dataset
data.isnull().sum()
```

```
Out[11]: type                0
fixed acidity              10
volatile acidity           8
citric acid                 3
residual sugar             2
chlorides                  2
free sulfur dioxide         0
total sulfur dioxide        0
density                    0
pH                          9
sulphates                  4
alcohol                     0
quality                     0
dtype: int64
```

```
In [12]: #check duplicated value  
data.duplicated().sum()
```

```
Out[12]: 1168
```


EDA of The Dataset

```
In [13]: #count the value of type  
data['type'].value_counts()
```

```
Out[13]: white    4898  
         red      1599  
         Name: type, dtype: int64
```


```
In [14]: #plot conutplot  
sns.countplot(data['type'])
```

```
Out[14]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe6c0137a60>
```



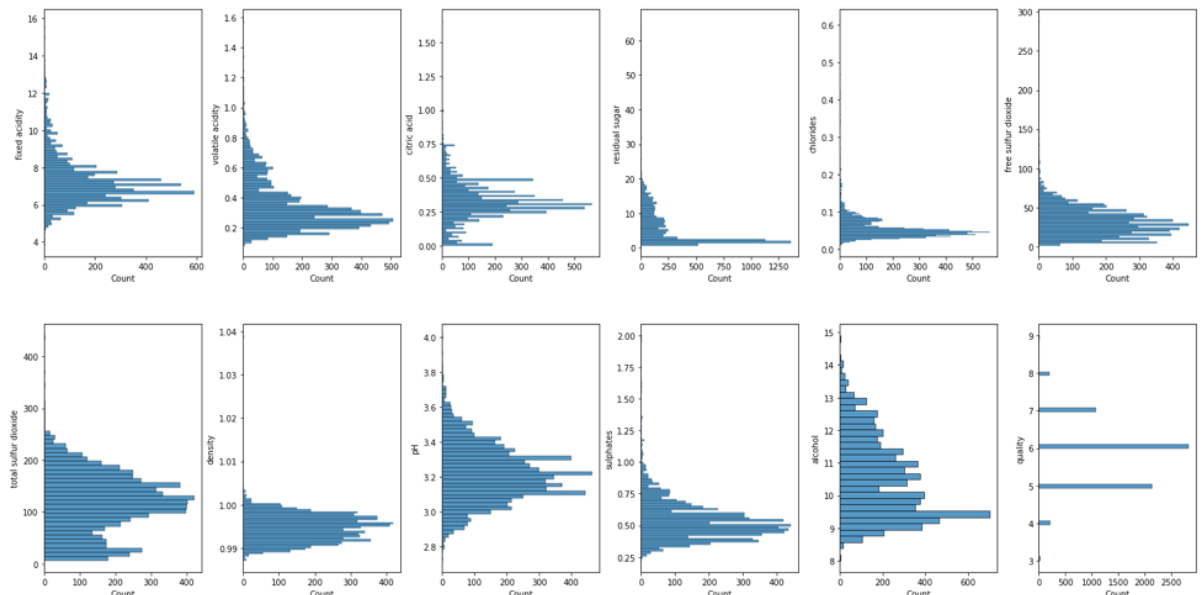
```
In [15]: #plot pie plot
data['type'].value_counts().plot(kind='pie', autopct='%.2f')
```

```
Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe6bf8c8e50>
```




```
In [16]: # create histplot plots
fig, ax = plt.subplots(ncols=6, nrows=2, figsize=(20,10))
index = 0
ax = ax.flatten()

for col, value in data.items():
    if col != 'type':
        sns.histplot(y=col, data=data, ax=ax[index])
        index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```




```
In [23]: #distplot for fixed acidity  
sns.distplot(data['fixed acidity'])
```

Out[23]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe6bec86520>




```
In [24]: #distplot for volatile acidity  
sns.distplot(data['volatile acidity'])
```

Out[24]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe6bee81ac0>



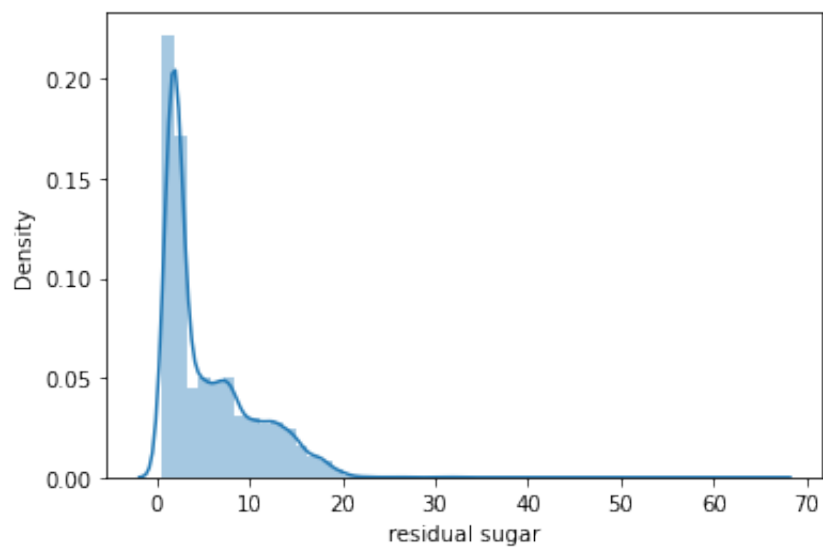
```
In [25]: #distplot for citric acid  
sns.distplot(data['citric acid'])
```

Out[25]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe6bf4aa970>




```
In [26]: #distplot for residual sugar  
sns.distplot(data['residual sugar'])
```

Out[26]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe6bf0d6fa0>





```
In [27]: #distplot for chlorides  
sns.distplot(data['chlorides'])
```

Out[27]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe6bf5531c0>




```
In [28]: #distplot for free sulfur dioxide  
sns.distplot(data['free sulfur dioxide'])
```

Out[28]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe6bf28f100>




```
In [29]: #distplot for total sulfur dioxide  
sns.distplot(data['total sulfur dioxide'])
```

Out[29]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe6bf22d580>




```
In [31]: #distplot for density  
sns.distplot(data['density'])
```

Out[31]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe6bfd46a90>




```
In [32]: #distplot for sulphates  
sns.distplot(data['sulphates'])
```

Out[32]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe6ba7148b0>




```
In [35]: #distplot for PH  
sns.distplot(data['pH'])
```

Out[35]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe6ba1dea90>




```
In [36]: #distplot for alcohol  
sns.distplot(data['alcohol'])
```

Out[36]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe6bec61bb0>



```
In [37]: #distplot for quality  
sns.distplot(data['quality'])
```

Out[37]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe6bed212b0>




```
In [38]: #count the value quality
data['quality'].value_counts()
```

```
Out[38]: 6    2836
         5    2138
         7    1079
         4     216
         8     193
         3      30
         9       5
         Name: quality, dtype: int64
```


```
In [39]: #countplot value quality
sns.countplot(data['quality'])
```

```
Out[39]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe6beca4160>
```



```
In [40]: #plot pie plot
data['quality'].value_counts().plot(kind='pie', autopct='%.2f')
```


```
Out[40]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe6bf7aefd0>
```



```
In [41]: # log transformation
data['free sulfur dioxide'] = np.log(1 + data['free sulfur dioxide'])
```

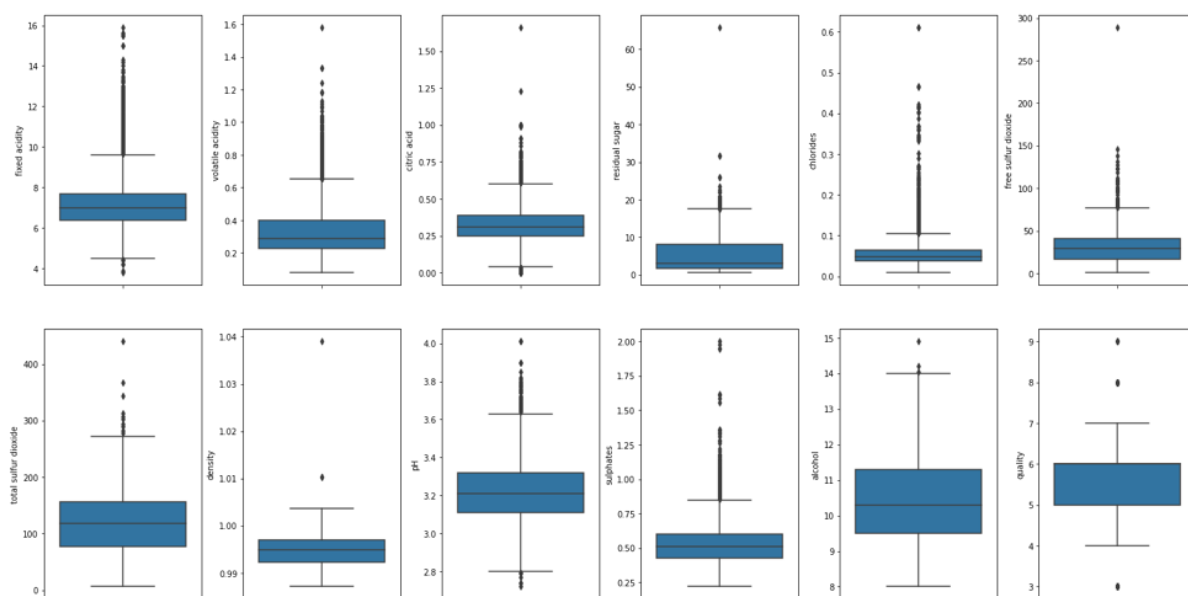
```
In [42]: sns.distplot(data['free sulfur dioxide'])
```

```
Out[42]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe6bc0dedc0>
```




```
In [18]: # create box plots
fig, ax = plt.subplots(ncols=6, nrows=2, figsize=(20,10))
index = 0
ax = ax.flatten()

for col, value in data.items():
    if col != 'type':
        sns.boxplot(y=col, data=data, ax=ax[index])
        index += 1
plt.tight_layout(pad=0.5, w_pad=0.7, h_pad=5.0)
```



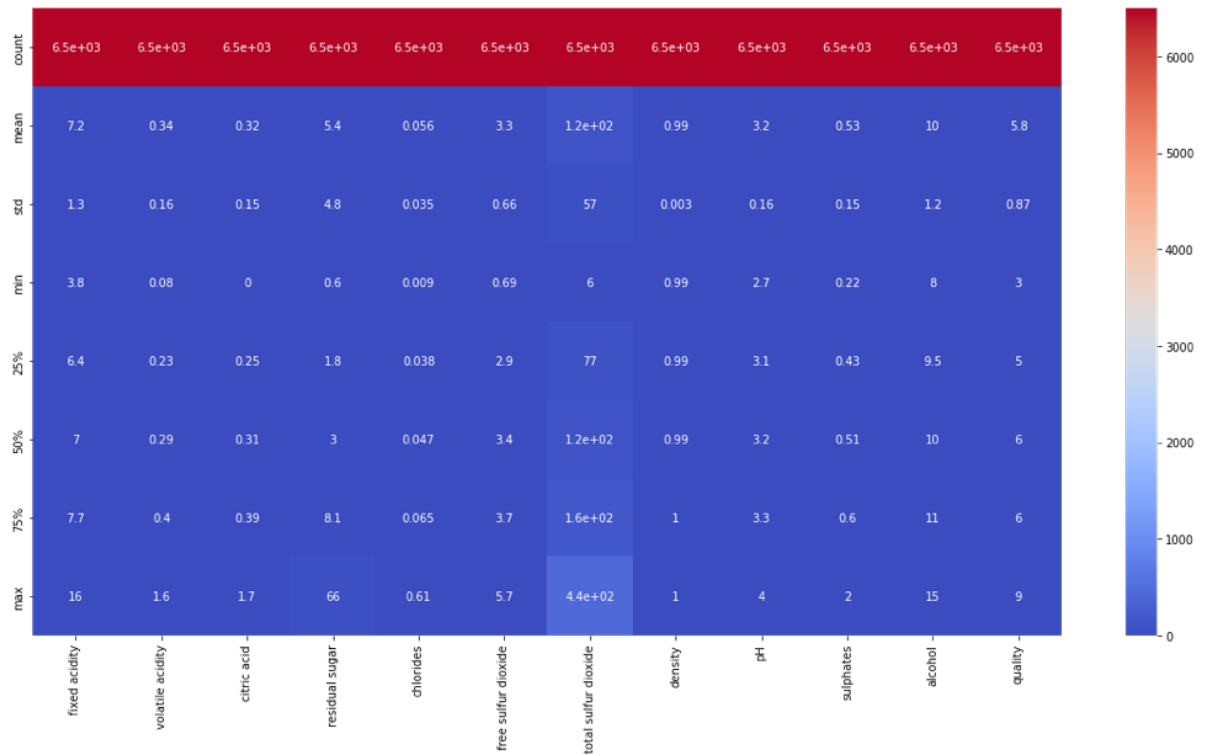
```
In [43]: #correaction of the dataset
corr = data.corr()
plt.figure(figsize=(20,10))
sns.heatmap(corr, annot=True, cmap='coolwarm')
```

Out[43]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe6bf1856a0>




```
In [46]: desc = data.describe()  
plt.figure(figsize=(20,10))  
sns.heatmap(desc, annot=True, cmap='coolwarm')
```

Out[46]: <matplotlib.axes._subplots.AxesSubplot at 0x7fe6bec22b50>




```
In [48]: #pairplot of the dataset
sns.pairplot(data,hue='quality')
```

```
Out[48]: <seaborn.axisgrid.PairGrid at 0x7fe6b64a2460>
```



```
In [55]: # fill the missing values
for col, value in data.items():
    if col != 'type':
        data[col] = data[col].fillna(data[col].mean())
```

```
In [56]: #after fill missing value check missing value
data.isnull().sum()
```

```
Out[56]: type                0
fixed acidity              0
volatile acidity           0
citric acid                0
residual sugar             0
chlorides                  0
free sulfur dioxide        0
total sulfur dioxide       0
density                    0
pH                         0
sulphates                  0
alcohol                    0
quality                    0
dtype: int64
```

```
In [57]: #splitting the dataset in X and Y
X = data.drop(columns=['type', 'quality'])
Y = data['quality']
```

```
In [58]: print(X)
print(Y)
```

	fixed acidity	volatile acidity	citric acid	residual sugar
0	7.0	0.270	0.36	20.7
1	6.3	0.300	0.34	1.6
2	8.1	0.280	0.40	6.9
3	7.2	0.230	0.32	8.5
4	7.2	0.230	0.32	8.5
...
6492	6.2	0.600	0.08	2.0
6493	5.9	0.550	0.10	2.2
6494	6.3	0.510	0.13	2.3
6495	5.9	0.645	0.12	2.0
6496	6.0	0.310	0.47	3.6
	free sulfur dioxide	total sulfur dioxide	density	pH
	chlorides \			su
	lphates \			

```

0          3.828641          170.0  1.00100  3.00  0
.450000
1          2.708050          132.0  0.99400  3.30  0
.490000
2          3.433987           97.0  0.99510  3.26  0
.440000
3          3.871201          186.0  0.99560  3.19  0
.400000
4          3.871201          186.0  0.99560  3.19  0
.400000
...          ...          ...          ...          ...
...
6492        3.496508          44.0  0.99490  3.45  0
.580000
6493        3.688879          51.0  0.99512  3.52  0
.531215
6494        3.401197          40.0  0.99574  3.42  0
.750000
6495        3.496508          44.0  0.99547  3.57  0
.710000
6496        2.944439          42.0  0.99549  3.39  0
.660000

```

```

      alcohol
0         8.8
1         9.5
2        10.1
3         9.9
4         9.9
...      ...
6492      10.5
6493      11.2
6494      11.0
6495      10.2
6496      11.0

```

[6497 rows x 11 columns]

```

0         6
1         6
2         6
3         6
4         6
...
6492      5
6493      6
6494      6
6495      5
6496      6

```

Name: quality, Length: 6497, dtype: int64

```
In [63]: Y.value_counts()
```

```
Out[63]: 6      2836
         5      2138
         7      1079
         4       216
         8       193
         3        30
         9         5
         Name: quality, dtype: int64
```

```
In [64]: from imblearn.over_sampling import SMOTE
oversample = SMOTE(k_neighbors=4)
# transform the dataset
X, Y = oversample.fit_resample(X, Y)
```

```
In [68]: Y.value_counts()
```

```
Out[68]: 6      2836
         5      2836
         7      2836
         8      2836
         4      2836
         3      2836
         9      2836
         Name: quality, dtype: int64
```

```
In [59]: #splitting the dataset in X_train and Y_train
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
```

```
In [65]: # split the data to train and test set
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,train_size=0.8
```

```
In [66]: #print X_train and Y_train
print(X_train.shape, X_test.shape, Y_train.shape, Y_test.shape)

(16874, 11) (2978, 11) (16874,) (2978,)
```

```
In [67]: #usingg standardscaler
scaler = StandardScaler()
X_train =scaler.fit_transform(X_train)
```

Model Training

```
In [69]: # classify function
from sklearn.model_selection import cross_val_score, train_test_split
def classify(model, X_train, Y_train):
    # train the model
    model.fit(X_train, Y_train)
    print("Accuracy:", model.score(X_test, Y_test) * 100)

    # cross-validation
    score = cross_val_score(model, X_train, Y_train, cv=5)
    print("CV Score:", np.mean(score)*100)
```

```
In [70]: #using LogisticRegression
from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
classify(model, X_train, Y_train)
```

Accuracy: 8.059100067159166
CV Score: 52.275715383433216

```
In [71]: from sklearn.tree import DecisionTreeClassifier
model = DecisionTreeClassifier()
classify(model, X_train, Y_train)
```

Accuracy: 13.834788448623236
CV Score: 79.01506860743376

```
In [72]: from sklearn.ensemble import RandomForestClassifier
model = RandomForestClassifier()
classify(model, X_train, Y_train)
```

Accuracy: 15.547347212894561
CV Score: 87.37703308524885

```
In [73]: from sklearn.ensemble import ExtraTreesClassifier
model = ExtraTreesClassifier()
classify(model, X_train, Y_train)
```

Accuracy: 14.707857622565479
CV Score: 88.68078772311137

In []: